

Overview of regression techniques

Point count data analysis workshop 2025

Péter Sólymos

2025-11-12

Table of contents

Preamble	1
Covariates	2
Variable types	2
Data exploration	7
Variable transformations	14
Compound variables	14
Modeling	15
Poisson null model	18
Main effects	20
Non-linear effects	23
Categorical variables	24
Multiple main effects	27
Interactions	28
Different error distributions	32
Mixed models	35
Next	38

Preamble

```
suppressPackageStartupMessages({  
    library(dplyr)  
    library(ggplot2)
```

```
library(mefa4)
library(detect)
})
```

Covariates

Variables that co-vary with the response variable. Also called as independent variables, predictors.

Let's continue with the JOSM data set:

```
x <- detect::josm$surveys |>
  select(
    Longitude,
    Latitude,
    WindStart,
    TSSR,
    DAY,
    Open,
    Water,
    Decid,
    OpenWet,
    Conif,
    ConifWet,
    Agr,
    UrbInd,
    SoftLin,
    Roads
  )
```

STOP AND EXPLAIN EACH VARIABLE!

Variable types

What type of variables we have? You can use the `str()` function to reveal the structure of R objects:

```
str(x)
```

```
'data.frame': 4569 obs. of 15 variables:
$ Longitude: num -113 -113 -113 -113 -113 ...
$ Latitude : num 55.2 55.2 55.2 55.2 55.2 ...
$ WindStart: int 0 0 0 0 0 1 2 0 0 0 ...
$ TSSR     : num 0.0132 0.0666 0.0125 0.041 0.1263 ...
$ DAY      : num 0.471 0.471 0.471 0.471 0.471 ...
$ Open     : num 0 0 0 0 0 0 0 0 0 0 ...
$ Water    : num 0.02055 0.00752 0.00752 0.03284 0.00416 ...
$ Decid   : num 0.8569 0.0443 0.0443 0.8823 0.8139 ...
$ OpenWet  : num 0.00315 0.58355 0.58355 0.06961 0.11164 ...
$ Conif   : num 0.0249 0 0 0 0 ...
$ ConifWet: num 0.0368 0.3456 0.3456 0 0.0018 ...
$ Agr     : num 0 0 0 0 0 0 0 0 0 0 ...
$ UrbInd  : num 0 0 0 0.01056 0.00839 ...
$ SoftLin : num 0.00521 0.01811 0.01811 0.00141 0.01349 ...
$ Roads   : num 0.052461 0.000913 0.000913 0.003307 0.046578 ...
```

We see that these variables are all continuous.

However, `WindStart` has very few distinct values, so we could treat it as ordinal (ordered factor):

```
table(x$WindStart)
```

	0	1	2	3	4	5	6
2471	926	844	287	31	5	3	

```
x$WindOrd <- as.ordered(x$WindStart)
str(x$WindOrd)
```

```
Ord.factor w/ 7 levels "0" <"1" <"2" <"3" <...: 1 1 1 1 1 2 3 1 1 1 ...
```

```
levels(x$WindOrd)
```

```
[1] "0" "1" "2" "3" "4" "5" "6"
```

Sometimes variables are binary. In R these can be logical (TRUE/FALSE) or coded as 0/1. Often we make such variables by discretizing other continuous or ordinal variables. E.g. We can create a binary wind variable:

```
x$Wind01 <- ifelse(x$WindStart > 0, 1, 0)
table(x$WindStart, x$Wind01)
```


	0	1
0	2471	0
1	0	926
2	0	844
3	0	287
4	0	31
5	0	5
6	0	3

Some categorical variables depend on some kind of classification, for example we can cut a continuous variable into bins:

```
x$DecidCut <- cut(x$Decid, seq(0, 1, 0.2), include.lowest = TRUE)
table(x$DecidCut)
```

[0,0.2]	(0.2,0.4]	(0.4,0.6]	(0.6,0.8]	(0.8,1]
1592	878	857	638	604

```
boxplot(Decid ~ DecidCut, x)
```



We can also inspect the land cover proportions that add up to 1 for each row.

```

## define column names
cn <- c(
  "Open", "Water", "Agr", "UrbInd", "SoftLin", "Roads", "Decid",
  "OpenWet", "Conif", "ConifWet"
)
## these sum to 1
summary(rowSums(x[, cn]))

```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1	1	1	1	1	1

The `find_max()` function finds the maximum value in each row, the output contains the value and the column where it was found, we can turn that into the dominant land cover type encoded in HAB:


```

h <- find_max(x[, cn])
head(h)

      index      value
CL10102 Decid 0.8569106
CL10106 OpenWet 0.5835472
CL10108 OpenWet 0.5835472
CL10109 Decid 0.8822829
CL10111 Decid 0.8139365
CL10112 Decid 0.8139365

```

```
hist(h$value)
```



```
table(h$index)
```

Open	Water	Agr	UrbInd	SoftLin	Roads	Decid	OpenWet
12	10	4	14	0	2	2084	160
Conif	ConifWet						
745	1538						

```
x$HAB <- droplevels(h$index) # drop empty levels  
x$DEC <- ifelse(x$HAB == "Decid", 1, 0)
```

Other types of categorical variables are truly discrete, like observer, where there are no underlying continuous data:

```
table(detect::josm$surveys$ObserverID)
```

2	14	15	19	22	26	27	28	32	41	48	57	59	63	64	65	69	75	82	86
186	223	88	267	7	191	307	227	231	140	170	126	222	240	127	154	302	127	173	280
87	89	93	98																
166	99	244	272																

When the number of categories increase and approach the sample size, we can consider treating these variables as random effects. E.g. the SurveyArea variable that has 271 levels.


Data exploration

We should inspect each variable that we want to use as a covariate. Here are some of the most important functions:

```
summary(x$Decid) # check mean, range, missing values
```

```
Min. 1st Qu. Median Mean 3rd Qu. Max.  
0.00000 0.09852 0.36132 0.38717 0.63232 0.99891
```

```
hist(x$Decid) # check skew and outliers
```



A nice way of getting all of the above nicely formatted is to use the `skimr` package:

```
skimr::skim(x)
```

Table 1: Data summary

Name	x
Number of rows	4569
Number of columns	20
Column type frequency:	
factor	3

numeric

17

Group variables	None
-----------------	------

Variable type: factor


	skim_variablen_missing	complete_rate	ordered	n_unique	top_counts
WindOrd	2	1	TRUE	7	0: 2471, 1: 926, 2: 844, 3: 287
DecidCut	0	1	FALSE	5	[0,: 1592, (0.: 878, (0.: 857, (0.: 638
HAB	0	1	FALSE	9	Dec: 2084, Con: 1538, Con: 745, Ope: 160

Variable type: numeric


	skim_variable_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
Longitude	0	1	-	1.96	-	-	-	-	-	-
			113.20	117.09	114.55	112.75	111.49	110.06		
Latitude	0	1	56.20	0.74	54.58	55.58	56.15	56.81	57.71	
WindStart	2	1	0.80	1.02	0.00	0.00	0.00	2.00	6.00	
TSSR	0	1	0.10	0.06	-0.03	0.05	0.10	0.16	0.24	
DAY	0	1	0.45	0.03	0.39	0.42	0.45	0.47	0.50	
Open	0	1	0.01	0.04	0.00	0.00	0.00	0.00	0.65	
Water	0	1	0.01	0.04	0.00	0.00	0.00	0.00	0.81	
Decid	0	1	0.39	0.30	0.00	0.10	0.36	0.63	1.00	
OpenWet	0	1	0.07	0.12	0.00	0.00	0.02	0.08	0.87	
Conif	0	1	0.18	0.22	0.00	0.01	0.08	0.27	1.00	
ConifWet	0	1	0.30	0.29	0.00	0.04	0.20	0.50	1.00	
Agr	0	1	0.00	0.02	0.00	0.00	0.00	0.00	0.79	
UrbInd	0	1	0.01	0.05	0.00	0.00	0.00	0.01	1.00	
SoftLin	0	1	0.02	0.02	0.00	0.01	0.01	0.03	0.25	
Roads	0	1	0.01	0.02	0.00	0.00	0.00	0.01	0.27	
Wind01	2	1	0.46	0.50	0.00	0.00	0.00	1.00	1.00	
DEC	0	1	0.46	0.50	0.00	0.00	0.00	1.00	1.00	

To explore relationships between variables, make scatter and box plots:


```
x |> ggplot(aes(x = Decid, y = Conif)) +  
  geom_point()
```




```
x |> ggplot(aes(x = Longitude, y = Latitude)) +  
  geom_point()
```



```
x |> ggplot(aes(x = DAY, y = TSSR)) +  
  geom_point()
```




```
x |> ggplot(aes(x = HAB, y = Decid)) +  
  geom_boxplot()
```



We can present 3 variables as color scatter or bubble plots:

```
x |> ggplot(aes(x = Decid, y = ConifWet, col = HAB)) +  
  geom_point()
```




Multivariate exploration include checking correlations:

```
round(cor(x[, cn]), 3)
```


	Open	Water	Agr	UrbInd	SoftLin	Roads	Decid	OpenWet	Conif
Open	1.000	-0.002	0.004	-0.007	0.006	0.011	-0.038	0.081	-0.019
Water	-0.002	1.000	0.023	0.007	-0.026	0.048	-0.039	0.012	-0.063
Agr	0.004	0.023	1.000	-0.004	-0.027	0.036	0.008	-0.012	-0.040
UrbInd	-0.007	0.007	-0.004	1.000	0.178	0.174	-0.100	-0.060	-0.070
SoftLin	0.006	-0.026	-0.027	0.178	1.000	0.193	-0.185	0.001	-0.152
Roads	0.011	0.048	0.036	0.174	0.193	1.000	0.055	-0.063	-0.037
Decid	-0.038	-0.039	0.008	-0.100	-0.185	0.055	1.000	-0.197	-0.314
OpenWet	0.081	0.012	-0.012	-0.060	0.001	-0.063	-0.197	1.000	-0.173
Conif	-0.019	-0.063	-0.040	-0.070	-0.152	-0.037	-0.314	-0.173	1.000
ConifWet	-0.115	-0.068	-0.048	-0.008	0.187	-0.120	-0.674	-0.069	-0.330
	ConifWet								
Open		-0.115							
Water		-0.068							
Agr		-0.048							
UrbInd		-0.008							
SoftLin		0.187							
Roads		-0.120							
Decid		-0.674							
OpenWet		-0.069							

```
Conif      -0.330
ConifWet    1.000
```

```
corrplot::corrplot(cor(x[, cn]), "ellipse")
```



```
heatmap(as.matrix(x[, cn]))
```



Variable transformations

We have seen examples of these:

- indicator variables (0/1)
- discretization (cut)

Other transformations include:

- sqrt: to tame outliers (not suitable for negative values)
- log: we'll see many use cases later
- polynomials: nonlinear terms (x^2 , x^3 , etc.)
- centering: keeps the distribution but shifts the mean
- scaling: keeps the distribution but shifts the range (often used with centering)

Compound variables

We can reduce correlation by combining variables together when those are additive:

```
x$FOR <- x$Decid + x$Conif + x$ConifWet
x$HF <- x$Agr + x$UrbInd + x$Roads + x$SoftLin
x$WET <- x$OpenWet + x$ConifWet + x$Water
```


We can reduce colinearity by calculating variables relative to each other, like proportions or ratios (watch out for division by 0):

```
x$pDecid <- ifelse(x$FOR > 0, x$Decid / x$FOR, 0)
cor(x[, c("FOR", "pDecid")])
```

```
            FOR      pDecid
FOR    1.00000000 0.01124794
pDecid 0.01124794 1.00000000
```

We can also merge categories, e.g. based on their similarity:

```
heatmap(cor(x[, cn]))
```



Modeling

We have manipulated the covariates and the species counts. Let's put them together in the same data frame.

```
spp <- "OVEN" # change here if you want to use another species
detect::josm$species[spp, ]
```

	SpeciesID	SpeciesName	ScientificName
OVEN	OVEN	Ovenbird	Seiurus aurocapillus

```
y <- mefa4::Xtab(~ SiteID + SpeciesID, detect::josm$counts)[, spp, drop = FALSE]
x$Count <- y[rownames(x), ]
```

Let's check for missing values:

```
data.frame(missing = colSums(is.na(x))) |> filter(missing > 0)
```

	missing
WindStart	2
WindOrd	2
Wind01	2

There are only 2 rows with missing values, we could use `na.omit(x)` to drop these, or we can impute the values:

- pick a value randomly
- use the mean
- use the most common value (mode)
- Use other variables to predict the possible value

After inspecting the full data set, it turns out that site FT21204W experiences wind according to `WindEnd`:

```
sites <- rownames(x)[is.na(x$WindStart)]
detect::josm$surveys[detect::josm$surveys$SiteID %in% c("FT07424", "FT21204W"), ]
```

	SiteID	SurveyArea	Longitude	Latitude	Date	StationID	
FT07424	FT07424	FT074	-111.7273	55.18270	2012-06-28	FT07424-1	
FT21204W	FT21204W	FT212	-111.5415	57.15522	2014-06-30	FT21204W-1	
	ObserverID	TimeStart	VisitID	WindStart	PrecipStart	TempStart	
FT07424	57	8:20:00 AM	1	NA	0	NA	
FT21204W	32	8:30:00 AM	1	NA	NA	20.3	
	CloudStart	WindEnd	PrecipEnd	TempEnd	CloudEnd	TimeFin	Noise
FT07424	0	NA	0	NA	0	8:30:00 AM	0

```

FT21204W      NA      3      0    20.3      0 8:40:00 AM      1
OvernightRain      DateTime      SunRiseTime SunRiseFrac
FT07424      TRUE 2012-06-28 08:20:00 2012-06-28 04:49:30  0.2010437
FT21204W      FALSE 2014-06-30 08:30:00 2014-06-30 04:34:45  0.1907999
TSSR OrdinalDay      DAY Open      Water Agr      UrbInd
FT07424  0.1461785      179 0.4904110      0 0.0023198943  0 0.006284186
FT21204W  0.1633668      180 0.4931507      0 0.0006266884  0 0.082233995
SoftLin      Roads Decid OpenWet Conif ConifWet
FT07424  0.000000000 0.001462686 0.2005029 0.03073732 0.4908488 0.2678442
FT21204W  0.008925667 0.061308954 0.2785886 0.000000000 0.0000000 0.5683161

```

Let's set the value to 3 for site FT21204W, and the mode (0) for site FT07424:

```

x[sites, "WindStart"] <- c(0, 3)
x[sites, "WindOrd"] <- c("0", "3")
x[sites, "Wind01"] <- c(0, 1)

```

Any more NAs left?

```
any(is.na(x))
```

```
[1] FALSE
```

Let's inspect the response variable:

```
table(x$Count)
```

Count	Frequency
0	2492
1	881
2	654
3	365
4	134
5	30
6	13

The distribution looks somewhat 0 inflated. What are the possible reasons for that?

- conditions lead to the absence of the species
- detected counts are lower than the actual counts

The types of models used most often to model count data and the functions used to fit them:

- Poisson: `stats::glm()`
- Negative Binomial (higher or lower variance than Poisson): `MASS::glm.nb()`
- Zero-inflated Poisson (ZIP): `pscl::zeroinfl()`

- Zero-inflated Negative Binomial (ZINB): `:: pscl:::zeroinfl()`

Other approaches include additive models (`mgcv:::gam()`) and tree based methods (see the `gbm` and `xgboost` packages).

Let us start with the Poisson model.

Poisson null model

The null model states that the expected values of the count at all locations are identical: $E[Y_i] = \lambda$ ($i = 1, \dots, n$), where Y_i is a random variable that follows a Poisson distribution with mean λ : $(Y_i | \lambda) \sim \text{Poisson}(\lambda)$. The observation (y_i) is a realization of the random variables Y at site i , these observations are independent and identically distributed (i.i.d.), and we have n observations in total.

Saying the the distribution is Poisson is an assumption in itself. For example we assume that the variance equals the mean ($V(\mu) = \mu$).

```
mP0 <- glm(Count ~ 1, data = x, family = poisson)
```

The `family=poisson` specification implicitly assumes that we use a logarithmic link functions, that is to say that $\log(\lambda) = \beta_0$, or equivalently: $\lambda = e^{\beta_0}$. The mean of the observations equal the mean of the fitted values, as expected:

```
mean(mP0$y)
```

```
[1] 0.8859707
```

```
mean(fitted(mP0))
```

```
[1] 0.8859707
```

```
exp(coef(mP0))
```

```
(Intercept)
0.8859707
```

The logarithmic function is called the link function, its inverse, the exponential function is called the inverse link function. The model family has these conveniently stored for us:

```
mP0$family
```

```
Family: poisson  
Link function: log
```

```
mP0$family$linkfun
```

```
function (mu)  
log(mu)  
<environment: namespace:stats>
```

```
mP0$family$linkinv
```

```
function (eta)  
pmax(exp(eta), .Machine$double.eps)  
<environment: namespace:stats>
```

Inspect the summary

```
summary(mP0)
```

```
Call:  
glm(formula = Count ~ 1, family = poisson, data = x)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.12107	0.01572	-7.703	1.33e-14 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

```
Null deviance: 7447.9 on 4568 degrees of freedom  
Residual deviance: 7447.9 on 4568 degrees of freedom  
AIC: 12603
```

Number of Fisher Scoring iterations: 6

Notice that the residual deviance much higher than residual degrees of freedom. This indicates that our parametric model (Poisson error distribution, constant expected value) is not quite right. See if we can improve that somehow and explain more of the variation.

We can pick an error distribution that would fit the residuals around the constant expected value better (i.e. using random effects). But this way we would not learn about what is driving the variation in the counts. We would also have a really hard time predicting abundance of the species for unsurveyed locations. We would be right on average, but we wouldn't be able to tell how abundance varies along e.g. a disturbance gradient or with tree cover.

An alternative approach would be to find predictors that could explain the variation.

Main effects

We fit a parametric (Poisson) linear model using `Decid` as a predictor:

```
mP1 <- glm(Count ~ Decid, data = x, family = poisson)
mean(mP1$y)
```

```
[1] 0.8859707
```

```
mean(fitted(mP1))
```

```
[1] 0.8859707
```

```
coef(mP1)
```

	(Intercept)	Decid
	-1.158833	2.130040

Same as before, the mean of the observations equal the mean of the fitted values. But instead of only the intercept, now we have 2 coefficients estimated. Our linear predictor thus looks like: $\log(\lambda_i) = \beta_0 + \beta_1 x_{1i}$. This means that expected abundance is e^{β_0} where `Decid`=0, $e^{\beta_0}e^{\beta_1}$ where `Decid`=1, and $e^{\beta_0+\beta_1 x_1}$ in between.


The relationship can be visualized by plotting the fitted values against the predictor, or using the coefficients to make predictions using our formula:

```

## make a sequence between 0 and 1
dec <- seq(from = 0, to = 1, by = 0.01)
## predict lambda
lam <- exp(coef(mP1)[1] + coef(mP1)[2] * dec)

plot(fitted(mP1) ~ Decid, x, pch = 19, col = "grey") # fitted
lines(lam ~ dec, col = 2) # our predicted
rug(x$Decid) # observed x values

```



The model summary tells us that residuals are not quite right. The residual deviance is still higher than residual degrees of freedom (these should be close if the Poisson assumption holds, but it is much better than what we saw for the null model).

We also learned that the `Decid` effect is significant (meaning that the effect size is large compared to the standard error):

```
summary(mP1)
```

```

Call:
glm(formula = Count ~ Decid, family = poisson, data = x)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.15883    0.03507 -33.04   <2e-16 ***
Decid       2.13004    0.05359   39.75   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
(Dispersion parameter for poisson family taken to be 1)
```

```
Null deviance: 7447.9 on 4568 degrees of freedom
Residual deviance: 5760.2 on 4567 degrees of freedom
AIC: 10917
```

```
Number of Fisher Scoring iterations: 6
```

Note: we see a significant (<0.05) P -value for the intercept as well. It is totally meaningless. That P -value relates to the null hypothesis of the intercept (β_0) being 0. There is nothing special about that, it is like saying the average abundance is different from 1.

But when β_1 is significantly different from 0, it means that the main effect has non-negligible effect on the mean abundance.

We can compare this model to the null (constant, intercept-only) model:

```
AIC(mP0, mP1)
```

	df	AIC
mP0	1	12602.84
mP1	2	10917.16

```
BIC(mP0, mP1)
```

	df	BIC
mP0	1	12609.27
mP1	2	10930.02

```
MuMIn::model.sel(mP0, mP1)
```

Model selection table

	(Intercept)	Decid	df	logLik	AICc	delta	weight
mP1	-1.1590	2.13	2	-5456.581	10917.2	0.00	1
mP0	-0.1211		1	-6300.422	12602.8	1685.68	0

Models ranked by AICc(x)

AIC uses the negative log likelihood and the number of parameters as penalty. Smaller value indicate a model that is closer to the (unknowable) true model (caveat: this statement is true only asymptotically, i.e. it holds for large sample sizes). For small samples, we often use BIC (more penalty for complex models when sample size is small), or AICc (as in `MuMIn::model.sel()`).

Non-linear effects

We can use polynomial terms to capture non (log) linear effects:


```
mP12 <- glm(Count ~ Decid + I(Decid^2), data = x, family = poisson)
mP13 <- glm(Count ~ Decid + I(Decid^2) + I(Decid^3), data = x, family = poisson)
mP14 <- glm(Count ~ Decid + I(Decid^2) + I(Decid^3) + I(Decid^4), data = x, family = poisson)
MuMIn::model.sel(mP0, mP1, mP12, mP13, mP14)
```

```
Model selection table
  (Intercept) Decid Decid^2 Decid^3 Decid^4 df logLik AICc delta weight
mP14 -2.6470 16.610 -38.740 41.78 -16.47 5 -5232.629 10475.3 0.00 1
mP13 -2.3730 11.340 -16.250 8.06          4 -5243.979 10496.0 20.70 0
mP12 -1.9090 6.209 -3.927          3 -5286.517 10579.0 103.77 0
mP1  -1.1590 2.130          2 -5456.581 10917.2 441.89 0
mP0  -0.1211          1 -6300.422 12602.8 2127.57 0
Models ranked by AICc(x)
```

Not a surprise that the most complex model won, we had enough degrees of freedoms to spare.

```
xnew <- data.frame(Decid = seq(0, 1, 0.01))

pr <- cbind(
  predict(mP1, xnew, type = "response"),
  predict(mP12, xnew, type = "response"),
  predict(mP13, xnew, type = "response"),
  predict(mP14, xnew, type = "response")
)
matplot(xnew$Decid, pr,
  lty = 1, type = "l",
  xlab = "Decid", ylab = "E[Y]"
)
legend("topleft",
  lty = 1, col = 1:4, bty = "n",
  legend = c("Linear", "Quadratic", "Cubic", "Quartic")
)
```



Categorical variables

Categorical variables are expanded into a *model matrix* before parameter estimation. The model matrix usually contains indicator variables for each level (value 1 when factor value equals a particular label, 0 otherwise) except for the *reference category* (check `relevel` if you want to change the reference category).

The estimate for the reference category comes from the intercept, the rest of the estimates are relative to the reference category. In the log-linear model example this means a ratio.

```
head(model.matrix(~DEC, x))
```

	(Intercept)	DEC
CL10102	1	1
CL10106	1	0
CL10108	1	0
CL10109	1	1
CL10111	1	1
CL10112	1	1

```
mP2 <- glm(Count ~ DEC, data = x, family = poisson)
summary(mP2)
```

```
Call:
glm(formula = Count ~ DEC, family = poisson, data = x)
```

```

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.85106    0.03070 -27.72   <2e-16 ***
DEC          1.21104    0.03574  33.89   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for poisson family taken to be 1)

```

Null deviance: 7447.9 on 4568 degrees of freedom
Residual deviance: 6123.1 on 4567 degrees of freedom
AIC: 11280

```

Number of Fisher Scoring iterations: 6

```
coef(mP2)
```

	DEC
(Intercept)	-0.8510608
	1.2110412

The estimate for a non-deciduous landscape is e^{β_0} , and it is $e^{\beta_0}e^{\beta_1}$ for deciduous landscapes. (Of course such binary classification at the landscape (1 km^2) level doesn't really make sense.)

```
MuMIn::model.sel(mP1, mP2)
```

Model selection table								
	(Intrc)	Decid	DEC	df	logLik	AICc	delta	weight
mP1	-1.1590	2.13		2	-5456.581	10917.2	0.00	1
mP2	-0.8511		1.211	2	-5638.039	11280.1	362.92	0
Models ranked by AICc(x)								

Having estimates for each land cover type improves the model, but the model using continuous variable is still better

```
mP3 <- glm(Count ~ HAB, data = x, family = poisson)
summary(mP3)
```

Call:

```

glm(formula = Count ~ HAB, family = poisson, data = x)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.3863   0.5774  -2.401  0.0163 *
HABWater     1.0296   0.6901   1.492  0.1357
HABAgr       0.6931   0.9129   0.759  0.4477
HABUrbInd    0.1335   0.7638   0.175  0.8612
HABRoads    -10.9163  201.2853  -0.054  0.9567
HABDecid     1.7463   0.5776   3.023  0.0025 **
HABOpenWet    0.4220   0.5914   0.714  0.4755
HABConif      0.9214   0.5792   1.591  0.1117
HABConifWet   0.2942   0.5790   0.508  0.6114
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for poisson family taken to be 1)

```

Null deviance: 7447.9 on 4568 degrees of freedom
Residual deviance: 6023.2 on 4560 degrees of freedom
AIC: 11194

```

Number of Fisher Scoring iterations: 10

```
MuMIn::model.sel(mP1, mP2, mP3)
```

Model selection table										
	(Intrc)	Decid	DEC	HAB	df	logLik	AICc	delta	weight	
mP1	-1.1590		2.13		2	-5456.581	10917.2	0.00	1	
mP3	-1.3860				+	9	-5588.059	11194.2	276.99	0
mP2	-0.8511			1.211	2	-5638.039	11280.1	362.92	0	

Models ranked by AICc(x)

The prediction in this case would look like: $\log(\lambda_i) = \beta_0 + \sum_{j=1}^{k-1} \beta_j x_{ji}$, where we have k factor levels (and $k - 1$ indicator variables besides the intercept).

Here is a general way of calculating fitted values or making predictions based on the design matrix (**X**) and the coefficients (**b**) (column ordering in **X** must match the elements in **b**) given a parametric log-linear model **object** and data frame **df** (the code won't run as is, **object** is just a placeholder for your GLM model object):

```
b <- coef(object)
X <- model.matrix(formula(object), df)
exp(X %*% b)
```

Multiple main effects

We can add main effects by providing a new formula to a new model. Or we can use the `update()` function. The `. ~ .` means to keep both sides of the formula but add (+) or remove (-) terms from the original scope. Note: the `update` function can be used to update any function call, not just GLM models.

```
# mP4 <- glm(Count ~ Decid + ConifWet, data=x, family=poisson)
mP4 <- update(mP1, . ~ . + ConifWet)
summary(mP4)
```

```
Call:
glm(formula = Count ~ Decid + ConifWet, family = poisson, data = x)

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.69255   0.05547 -12.485  <2e-16 ***
Decid        1.61480   0.07169  22.525  <2e-16 ***
ConifWet     -0.98628   0.09906  -9.957  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 7447.9 on 4568 degrees of freedom
Residual deviance: 5659.7 on 4566 degrees of freedom
AIC: 10819

Number of Fisher Scoring iterations: 6
```

```
MuMIn::model.sel(mP0, mP1, mP4)
```

Model selection table								
(Intrc)	Decid	CnfWt	df	logLik	AICc	delta weight		
mP4	-0.6925	1.615	-0.9863	3	-5406.310	10818.6	0.00	1

```

mP1 -1.1590 2.130      2 -5456.581 10917.2   98.54      0
mP0 -0.1211           1 -6300.422 12602.8  1784.22      0
Models ranked by AICc(x)

```

Here are some functions that can automate model selection:

- `drop1`: evaluate which variable to drop to lower AIC the most
- `add1`: evaluate which variable to add to lower AIC the most
- `step`: perform forward/backward model selection using `add1/drop1`

Interactions

When we consider interactions between two variables (say x_1 and x_2), we refer to adding another variable to the model matrix that is a product of the two variables ($x_{12} = x_1x_2$):

```
head(model.matrix(~ x1 * x2, data.frame(x1 = 1:4, x2 = 10:7)))
```

```

(Intercept) x1 x2 x1:x2
1            1  1 10    10
2            1  2  9    18
3            1  3  8    24
4            1  4  7    28

```

Let's consider interaction between our two predictors from before:

```
mP5 <- glm(Count ~ Decid * ConifWet, data = x, family = poisson)
summary(mP5)
```

```

Call:
glm(formula = Count ~ Decid * ConifWet, family = poisson, data = x)

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -0.55157   0.05643 -9.775   <2e-16 ***
Decid        1.20486   0.07799 15.448   <2e-16 ***
ConifWet     -2.32015   0.14875 -15.598   <2e-16 ***
Decid:ConifWet 5.34691   0.35615 15.013   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
(Dispersion parameter for poisson family taken to be 1)
```

```
Null deviance: 7447.9 on 4568 degrees of freedom
Residual deviance: 5415.8 on 4565 degrees of freedom
AIC: 10577
```

```
Number of Fisher Scoring iterations: 6
```

```
MuMIn::model.sel(mP0, mP1, mP4, mP5)
```

```
Model selection table
```


	(Int)	Dcd	CnW	CnW:Dcd	df	logLik	AICc	delta	weight
mP5	-0.5516	1.205	-2.3200	5.347	4	-5284.389	10576.8	0.00	1
mP4	-0.6925	1.615	-0.9863		3	-5406.310	10818.6	241.84	0
mP1	-1.1590	2.130			2	-5456.581	10917.2	340.38	0
mP0	-0.1211				1	-6300.422	12602.8	2026.06	0

```
Models ranked by AICc(x)
```

The model with the interaction is best supported, but how do we make sense of this relationship? We can't easily visualize it in a single plot. We can either

1. fix all variables (at their mean/meadian) and see how the response is changing along a single variable: this is called a *conditional* effect (conditional on fixing other variables), this is what `visreg::visreg()` does
2. or plot the fitted values against the predictor variable (one at a time), this is called a *marginal* effect, and this is what `ResourceSelection::mep()` does

```
visreg::visreg(mP5, scale = "response", xvar = "ConifWet", by = "Decid")
```




```
ggplot(x, aes(x = Decid, y = fitted(mP5), col = ConifWet)) +
  geom_point() +
  geom_smooth() +
  theme_light()
```

```
`geom_smooth()` using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'
```

Warning: The following aesthetics were dropped during statistical transformation:
colour.

- i This can happen when ggplot fails to infer the correct grouping structure in the data.
- i Did you forget to specify a `group` aesthetic or to convert a numerical variable into a factor?



Final battle of Poisson models:

```
MuMIn::model.sel(mP0, mP1, mP12, mP13, mP14, mP2, mP3, mP4, mP5)
```

Model selection table

	(Int)	Dcd	Dcd^2	Dcd^3	Dcd^4	DEC HAB	CnW	CnW:Dcd	df	logLik
mP14	-2.6470	16.610	-38.740	41.78	-16.47				5	-5232.629
mP13	-2.3730	11.340	-16.250	8.06					4	-5243.979
mP5	-0.5516	1.205					-2.3200	5.347	4	-5284.389
mP12	-1.9090	6.209	-3.927						3	-5286.517
mP4	-0.6925	1.615					-0.9863		3	-5406.310
mP1	-1.1590	2.130							2	-5456.581
mP3	-1.3860							+	9	-5588.059
mP2	-0.8511					1.211			2	-5638.039
mP0	-0.1211								1	-6300.422
	AICc	delta	weight							
mP14	10475.3	0.00	1							
mP13	10496.0	20.70	0							
mP5	10576.8	101.52	0							
mP12	10579.0	103.77	0							
mP4	10818.6	343.35	0							
mP1	10917.2	441.89	0							
mP3	11194.2	718.89	0							
mP2	11280.1	804.81	0							

```
mP0 12602.8 2127.57      0
Models ranked by AICc(x)
```

Of course, the most complex model wins but the Chi-square test is still significant (indicating lack of fit). Let's try different error distribution.

Different error distributions

We will use the 2-variable model with interaction:

```
mP <- glm(Count ~ Decid * ConifWet, data = x, family = poisson)
```

Let us try the Negative Binomial distribution first. This distribution is related to Binomial experiments (number of trials required to get a fixed number of successes given a binomial probability). It can also be derived as a mixture of Poisson and Gamma distributions (see [Wikipedia](#)), which is a kind of hierarchical model. In this case, the Gamma distribution acts as an i.i.d. random effect for the intercept: $Y_i \sim Poisson(\lambda_i)$, $\lambda_i \sim Gamma(e^{\beta_0 + \beta_1 x_{1i}}, \gamma)$, where γ is the Gamma variance.

The Negative Binomial variance (using the parametrization common in R functions) is a function of the mean and the scale: $V(\mu) = \mu + \mu^2/\theta$.

```
mNB <- MASS::glm.nb(Count ~ Decid * ConifWet, data = x)
summary(mNB)
```

Call:

```
MASS::glm.nb(formula = Count ~ Decid * ConifWet, data = x, init.theta = 3.524756181,
link = log)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.58101	0.06291	-9.236	<2e-16 ***
Decid	1.23695	0.08919	13.868	<2e-16 ***
ConifWet	-2.36485	0.16040	-14.743	<2e-16 ***
Decid:ConifWet	5.70380	0.40132	14.212	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Negative Binomial(3.5248) family taken to be 1)

```

Null deviance: 6086.2 on 4568 degrees of freedom
Residual deviance: 4388.4 on 4565 degrees of freedom
AIC: 10464

```

```
Number of Fisher Scoring iterations: 1
```

```

Theta: 3.525
Std. Err.: 0.411

```

```
2 x log-likelihood: -10453.887
```

Next, we look at zero-inflated models. In this case, the mixture distribution is a Bernoulli distribution and a count distribution (Poisson or Negative Binomial, for example). The 0's can come from both the zero and the count distributions, whereas the >0 values can only come from the count distribution: $A_i \sim Bernoulli(\varphi)$, $Y_i \sim Poisson(A_i\lambda_i)$.

The zero part of the zero-inflated models are often parametrized as probability of zero ($1 - \varphi$), as in the `pscl:::zeroinfl` function:

```

## Zero-inflated Poisson
mZIP <- pscl:::zeroinfl(Count ~ Decid * ConifWet | 1, x, dist = "poisson")
summary(mZIP)

```

```

Call:
pscl:::zeroinfl(formula = Count ~ Decid * ConifWet | 1, data = x, dist = "poisson")

Pearson residuals:
    Min      1Q   Median      3Q      Max 
-1.2160 -0.7011 -0.3391  0.3766  8.9394 

Count model coefficients (poisson with log link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -0.31127   0.06503 -4.787  1.69e-06 ***
Decid        1.05891   0.08588 12.330  < 2e-16 ***
ConifWet     -2.45044   0.15622 -15.686  < 2e-16 ***
Decid:ConifWet 5.94422   0.38368 15.493  < 2e-16 ***

Zero-inflation model coefficients (binomial with logit link):
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -1.5989     0.1021 -15.65  <2e-16 *** 
---
```

```
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Number of iterations in BFGS optimization: 12
```

```
Log-likelihood: -5218 on 5 Df
```

```
## Zero-inflated Negative Binomial  
mZINB <- pscl::zeroinfl(Count ~ Decid * ConifWet | 1, x, dist = "negbin")  
summary(mZINB)
```

```
Call:
```

```
pscl::zeroinfl(formula = Count ~ Decid * ConifWet | 1, data = x, dist = "negbin")
```

```
Pearson residuals:
```

Min	1Q	Median	3Q	Max
-1.1874	-0.6902	-0.3375	0.3616	8.9455

```
Count model coefficients (negbin with log link):
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.37625	0.07312	-5.146	2.66e-07 ***
Decid	1.10981	0.09105	12.189	< 2e-16 ***
ConifWet	-2.43272	0.15872	-15.327	< 2e-16 ***
Decid:ConifWet	5.93085	0.39641	14.962	< 2e-16 ***
Log(theta)	2.61874	0.51590	5.076	3.85e-07 ***

```
Zero-inflation model coefficients (binomial with logit link):
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.8513	0.1917	-9.655	<2e-16 ***

```
---
```

```
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Theta = 13.7184
```

```
Number of iterations in BFGS optimization: 22
```

```
Log-likelihood: -5215 on 6 Df
```

Now we compare the four different parametric models:

```
AIC(mP, mNB, mZIP, mZINB)
```

	df	AIC
mP	4	10576.78

```
mNB      5 10463.89
mZIP     5 10445.40
mZINB    6 10442.56
```

```
MuMIn::model.sel(mP, mNB, mZIP, mZINB)
```

```
Model selection table
  (Int)      CnW     Dcd CnW:Dcd cnt_(Int) cnt_CnW cnt_Dcd cnt_CnW:Dcd
mZINB                      -0.3762  -2.433   1.110          +
mZIP                      -0.3113  -2.450   1.059          +
mNB      -0.5810 -2.365  1.237   5.704
mP       -0.5516 -2.320  1.205   5.347
  zer_(Int)      family    class init.theta link dist df    logLik
mZINB     -1.851      b(lgt) zeroinfl                   ngbn  6 -5215.281
mZIP      -1.599      b(lgt) zeroinfl                   pssn  5 -5217.698
mNB        NB(3.5248,log) negbin      3.52   log      5 -5226.943
mP         p(log)      glm                    4 -5284.389
  AICc  delta weight
mZINB 10442.6  0.00  0.804
mZIP   10445.4  2.83  0.196
mNB    10463.9 21.32  0.000
mP     10576.8 134.21 0.000
Abbreviations:
family: b(lgt) = 'binomial(logit)',
          NB(3.5248,log) = 'Negative Binomial(3.5248,log)',
          p(log) = 'poisson(log)'
dist:   ngbn = 'negbin', pssn = 'poisson'
Models ranked by AICc(x)
```

Our best model is the ZINB. The probability of observing a zero as part of the zero distribution is back transformed from the zero coefficient using the inverse logit function:

```
unname(plogis(coef(mZIP, "zero"))) # P of 0 (not 1!)
```

```
[1] 0.1681348
```

Mixed models

It is also common practice to consider generalized linear mixed models (GLMMs) for count data. These mixed models are usually considered as Poisson-Lognormal mixtures. The simplest, so

called i.i.d., case is similar to the Negative Binomial, but instead of Gamma, we have Lognormal distribution: $Y_i \sim Poisson(\lambda_i)$, $\log(\lambda_i) = \beta_0 + \beta_1 x_{1i} + \epsilon_i$, $\epsilon_i \sim Normal(0, \sigma^2)$, where σ^2 is the Lognormal variance on the log scale.

We can use the `lme4::glmer` function: use `SiteID` as random effect (we have exactly n random effects).

```
mPLN1 <- lme4::glmer(Count ~ Decid * ConifWet + (1 | SiteID),
  data = data.frame(SiteID = rownames(x), x), family = poisson
)
summary(mPLN1)
```

```
Generalized linear mixed model fit by maximum likelihood (Laplace
Approximation) [glmerMod]
Family: poisson  ( log )
Formula: Count ~ Decid * ConifWet + (1 | SiteID)
Data: data.frame(SiteID = rownames(x), x)

AIC      BIC      logLik deviance df.resid
10445.9 10478.0  -5218.0   10435.9      4564

Scaled residuals:
    Min      1Q  Median      3Q     Max 
-1.1475 -0.6291 -0.2855  0.4156  5.4234

Random effects:
 Groups Name        Variance Std.Dev.
 SiteID (Intercept) 0.2988   0.5466 
Number of obs: 4569, groups: SiteID, 4569

Fixed effects:
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -0.74536   0.06743 -11.05   <2e-16 ***
Decid       1.27742   0.09201  13.88   <2e-16 ***
ConifWet    -2.34745   0.16248 -14.45   <2e-16 ***
Decid:ConifWet 5.63872   0.41229  13.68   <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:
          (Intr) Decid  ConifWt
Decid     -0.895
ConifWet   -0.621  0.644
```

```
Decid:CnfWt  0.176 -0.379 -0.700
```

Note: the number of unknowns we have to somehow estimate is now more than the number of observations we have. How is that possible?

Alternatively, we can use `SurveyArea` as a grouping variable. We have now $m < n$ random effects, and survey areas can be seen as larger landscapes within which the sites are clustered: $Y_{ij} \sim Poisson(\lambda_{ij})$, $\log(\lambda_{ij}) = \beta_0 + \beta_1 x_{1ij} + \epsilon_i$, $\epsilon_i \sim Normal(0, \sigma^2)$. The index i ($i = 1, \dots, m$) defines the cluster (survey area), the j ($j = 1, \dots, n_i$) defines the sites within survey area i ($n = \sum_{i=1}^m n_i$).

```
mPLN2 <- lme4::glmer(Count ~ Decid * ConifWet + (1 | SurveyArea),
  data = data.frame(SurveyArea = detect::josm$surveys$SurveyArea, x),
  family = poisson
)
summary(mPLN2)
```

```
Generalized linear mixed model fit by maximum likelihood (Laplace
Approximation) [glmerMod]
Family: poisson  ( log )
Formula: Count ~ Decid * ConifWet + (1 | SurveyArea)
Data: data.frame(SurveyArea = detect::josm$surveys$SurveyArea, x)

AIC      BIC      logLik deviance df.resid
10047.2  10079.3  -5018.6   10037.2     4564

Scaled residuals:
    Min      1Q      Median      3Q      Max 
-1.7379 -0.6434 -0.3200  0.3580  6.4556 

Random effects:
 Groups      Name        Variance Std.Dev. 
SurveyArea (Intercept) 0.2934   0.5417  
Number of obs: 4569, groups: SurveyArea, 271

Fixed effects:
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -0.73773   0.07807  -9.45   <2e-16 ***
Decid       1.19195   0.09820   12.14   <2e-16 ***
ConifWet    -2.32770   0.16824  -13.84   <2e-16 ***
Decid:ConifWet 5.53058   0.39698   13.93   <2e-16 ***
---

```

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:
  (Intr) Decid  ConfWt
Decid     -0.808
ConifWet   -0.609  0.628
Decid:CnfWt  0.162 -0.325 -0.670

```

In the battle of distributions (keeping the linear predictor part the same) the clustered GLMM was best supported:

```

tmp <- AIC(mP, mNB, mZIP, mZINB, mPLN1, mPLN2)
tmp$delta_AIC <- tmp$AIC - min(tmp$AIC)
tmp[order(tmp$AIC), ]

```

	df	AIC	delta_AIC
mPLN2	5	10047.19	0.0000
mZINB	6	10442.56	395.3703
mZIP	5	10445.40	398.2047
mPLN1	5	10445.91	398.7171
mNB	5	10463.89	416.6952
mP	4	10576.78	529.5872

Next

Naïve estimates of occupancy and abundance